

UNIVERSITA' DEGLI STUDI DELL'INSUBRIA
DIPARTIMENTO DI SCIENZE TEORICHE E
APPLICATE (DiSTA)
CORSO DI LAUREA MAGISTRALE IN INFORMATICA

INTELLIGENT SYSTEM PROJECT

RELAZIONE DI:
Vladi VALSECCHI matr 730030
Gianluca MOLTENI matr 730113

ANNO ACCADEMICO 2020/2021

INTRODUCTION

In this project we will proceed to the classification of two datasets, using two classification models and varying their metrics. We will then carry out an analysis of the results obtained for the single algorithm and an overall analysis of the performances.

Our goal is to understand which classification algorithm performs better on the inputs considered. The project includes a main file and a function file made on Jupyter Notebook in Python language (.ipynb format) using a Python 3.6 environment generated in Anaconda in which we implemented methods for choosing the features to be used for the classification and subsequently the actual classification.

PART A

What are supervised methods?

Supervised learning is a type of machine learning. Its goal is to classify data based on samples provided for reference. Each sample is associated with a label (represents the class). Consequently, the expression <pattern-class label> is used to indicate the association.

The characteristic of supervised methods is the separation of the data provided in input into two subsets called training set and test set. The first is the set used for learning and creating the model, the second is used for validating the model.

A mapping function is used to associate the input data with the most suitable class.

Supervised methods can be divided in two categories:

- Classification: is the process of prediction of the class based on the provided data. Classes are called targets/labels or categories;
- Regression: provides us with a simple linear relationship of two or more variables within a dataset.

The classification must be performed with a consistent set of data. This is because the performance varies according to the goodness of the training data provided in input. Another factor affecting performance is the amount of data provided. In this case, may happen a situation of:

- Overfitting - when there are many parameters that must be taken in account. The result is a high precision but fails in the generalization process;
- Underfitting - few parameters are provided during the training phase, in which the classifiers suffer from an excessive discrepancy.

PART B

Introduction

For the realization of the project, we relied on two supervised classification models.

With classification we mean the processes of prediction of data classes of points. The classes are also called target/labels or categories. Classification models require a set of input variables (X) and a mapping function (f) which associate each input with an output. The result (Y) is a discrete variable. Given an input $x \in X$ and W classes $\omega_1, \dots, \omega_W$ we must find W decision functions $d_1(x), \dots, d_W(x)$ with the property that a given x belongs to the class ω_i iff $d_i(x) > d_j(x)$ with $j=1, \dots, W$ and $j \neq i$. The function f is also called the discriminant function and it helps us to understand where two classes are separated. Such a decision boundary between two classes (ω_i, ω_j) is given by the value of x for which $d_i(x) = d_j(x)$.

The difference between two classes is given by:

$$d_{ij}(x) = d_i(x) - d_j(x) = 0$$

$d_{ij}(x) > 0 \rightarrow x$ belongs to ω_i

$d_{ij}(x) < 0 \rightarrow x$ belongs to ω_j

The supervised classifiers chosen are Minimum Distance Classifier (MDC) and K-Nearest Neighbor (KNN).

MINIMUM DISTANCE CLASSIFIER (MDC)

The Minimum Distance Classifier or MDC is a supervised classification algorithm that associates a given x to the class based on the minimum distance. The classifier makes use of Minkowski distances:

$$d(x, y) = (\sum_{i=1, N} (x_i - y_i)^\lambda)^{1/\lambda}$$

where λ is an integer and determines the metric:

- $\lambda=2$: Euclidean distance;
- $\lambda=1$: Cityblock/Manhattan distance;
- $\lambda \rightarrow \infty$: Chessboard;

The process includes the computation of a prototype vector for each class. An unknown pattern is compared with the prototype and assigned to the nearest class.

ADVANTAGES	DISADVANTAGES
Low demand for training data	Statistical information are ignored
Heuristic rule	Works well when the distance between means is large
Easy to design	
Low computational complexity	

K - NEAREST NEIGHBOR

The K Nearest Neighbor or KNN is a supervised classification algorithm that associates a given x to the class based on the quantity of elements of a given class in a given range. KNN is also based on Minkowski distances and the most used for this algorithm is the Euclidean distance. The classifier does not require a separate training phase and does not make statistical assumptions on the distribution of classes. KNN is a lazy learner because it does not learn a discriminative function from the training data but memorizes the training dataset instead.

Given W classes and N supervised training samples. To classify a pattern x , the distance is computed (using a predefined metrics) between the pattern x and k nearest samples. The class that has the most elements will be the class associated with the pattern.

ADVANTAGES	DISADVANTAGES
Easy to design and understand	Computational time at test time
No time to train	

PART C

The datasets used for the experiment are:

- Iris;
- Wine.

Both datasets come from the *sklearn.datasets* library.

The newly imported datasets are in the form of an array and were subsequently converted into dataframes using the *pandas* library, in order to better manage the data.

Dataset analysis

Iris


The dataset presents 150 instances divided into 3 classes Iris-Setosa, Iris-Versicolor and Iris-Virginica. Each class comprises 50 instances. The attributes that characterize the dataset are 4:

- sepal length in cm;
- sepal width in cm;
- petal length in cm;
- petal width in cm.

The data presents the following characteristics:

Feature	Min	Max	Mean	SD
<i>sepal length</i>	4.3	7.9	5.84	0.83
<i>sepal width</i>	2.0	4.4	3.05	0.43
<i>petal length</i>	1.0	6.9	3.76	1.76
<i>petal width</i>	0.1	2.5	1.20	0.76

We use a correlation matrix to determine the best correlation between the features. The greater the color intensity, the greater the relationship between features



Wine


The dataset presents 178 instances divided into 3 classes class-0, class-1 and class-2. The classes contain 59, 71 and 48 instances respectively. The attributes that characterize the dataset are 13:

- alcohol;
- malic_acid;
- ash;
- alcalinity_of_ash;
- magnesium;
- total_phenols;
- flavonoids;
- nonflavonoid_phenols;
- proanthocyanis;
- color_intensity;
- hue;
- od280/od315_of_diluted_wines;
- proline

The data presents the following characteristics:

Feature	Min	Max	Mean	SD
<i>alcohol</i>	11.0	14.8	13.0	0.8
<i>malic_acid</i>	0.74	5.80	2.34	1.12
<i>ash</i>	1.36	3.23	2.36	0.27
<i>alcalinity_of_ash</i>	1.06	30.0	19.5	3.3
<i>magnesium</i>	70.0	162.0	99.7	14.3
<i>total_phenols</i>	0.98	3.88	2.29	0.63
<i>flavonoids</i>	0.34	5.08	2.03	1.00
<i>nonflavonoid_phenols</i>	0.13	0.66	0.36	0.12
<i>proanthocyanis</i>	0.41	3.58	1.59	0.57
<i>color_intensity</i>	1.3	13.0	5.1	2.3
<i>hue</i>	0.48	1.71	0.96	0.23
<i>od280/od315_of_diluted_wines</i>	1.27	4.00	2.61	0.71
<i>proline</i>	278	1680	746	315

Again, we used a correlation matrix to analyze the relationship between features, and the results are represented below:



FEATURE SELECTION

For the feature selection operation we relied on the SelectKBest feature selector, coming from the *sklearn.feature_selection.SelectKBest* library. The selector uses chi2, so it computes chi-squared stats between each non-negative feature and class, and it is used to select the n_features (n_features=2) features with the highest values for the test chi-squared statistic.

K Best Iris

Analyzing the dataset and applying the feature selection method, we have:

Feature	Score	Pval
sepal length	10.817821	4.476515e-03
sepal width	3.710728	1.563960e-01
petal length	116.312613	5.533972e-26
petal width	67.048360	2.758250e-15

So, the k best features for Iris based on score are: **petal length, petal width.**

Kbest Wine

The same approach is applied on Wine dataset

Feature	Score	Pval
alcohol	5.445499	6.569389e-02
malic_acid	28.068605	8.034890e-07
ash	0.743381	6.895678e-01
alcalinity_of_ash	29.383695	4.163050e-07
magnesium	45.026381	1.669728e-10
total_phenols	15.623076	4.050346e-04
flavonoids	63.334308	1.766565e-14
nonflavonoid_phenols	1.815485	4.034340e-01
proanthocyanis	9.368283	9.240664e-03
color_intensity	109.016647	2.124887e-24
hue	5.182540	7.492483e-02
od280/od315_of_diluted_wines	23.389883	8.335878e-06
proline	16540.067145	0.000000e+00

The k best features for Wine based on score are: **proline, color_intensity**.

CLASSIFIERS UTILIZED

The models chosen for the experiments on the two mentioned datasets are the supervised classifiers:

- Minimum Distance Classifier;
- K-Nearest Neighbor.

MDC

For the use of this classifier, we have divided the training set and test set with a proportion of 70% and 30%. The number of resampling that is performed is also set to 40.

Metrics

Within the classifier class, you can set the metric for the distance calculation. The default metric is the Euclidean distance. Besides this metric we also used Manhattan distance and Chessboard (Chebyshev) distance.

KNN

Within the project we have implemented a second algorithm for the classification of patterns: the K-Nearest Neighbor.

K-Nearest Neighbor makes use of Minkowski's metrics, but in this case, we only used the default metric, which is the Euclidean metric.

At the beginning we divided the dataset into training set and test set, with a proportion of 70% and 30%.

We applied KNN with k increasing from 1 to 40 inclusive. We thus obtain <k - accuracy> pairs and we have analyzed a general case (k=5), the worst case and the best case.

This is implemented to see how the classifier performs in various scenarios.

Classifiers Evaluation

For the evaluation of the performance of the MDC with the three metrics used, we relied on accuracy. Accuracy is the number of correctly classified data instances out of the total number of data instances.

$$accuracy = \frac{true\ negative + true\ positive}{true\ negative + true\ positive + false\ positive + false\ negative}$$

Regarding the KNN we made use of the *confusion_matrix* and *classification_report* functions. Classification report provided us with the following parameters:

- Precision, is the ability of the classifier not to label as positive a sample that is negative.

$$precision = \frac{true\ positive}{true\ positive + false\ positive}$$

- Recall, indicates the percentage of true positives and the value;

$$recall = \frac{true\ positive}{(true\ positive + false\ negative)}$$

- F1-Score, is the harmonic mean between precision and recall and is a better measure of precision. The value fluctuates between 0 and 1;

$$f1 - score = 2 * \frac{Precision * Recall}{Precision + Recall}$$


- Support, is the number of occurrences of each class in y_test;
- Accuracy;
- Macro-AVG, calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.
- Weighted-AVG, calculate metrics for each label, and find their average weighted by support (the number of true instances for each label).

[[10 1 0] [1 15 0] [0 3 6]]					
	precision	recall	f1-score	support	
0	0.91	0.91	0.91	11	
1	0.79	0.94	0.86	16	
2	1.00	0.67	0.80	9	
accuracy			0.86	36	
macro avg	0.90	0.84	0.86	36	
weighted avg	0.88	0.86	0.86	36	


In order to compare the performance of MDC and KNN with the same accuracy metric, for the KNN we have chosen to rely exclusively on the accuracy metric.

The data that will be commented in the next sections are the results obtained from our experiment and may vary with new executions of the code. The performance evaluation of the two classifiers is based on these data.

MDC Euclidean Performance EVALUATION



Iris dataframe




Wine dataframe


As can be seen from the boxplots, the MDC with Euclidean metric, performs better in terms of accuracy, in the test phase, on the Iris dataframe, with an average accuracy of 96.33%, minimum outlier value of 91.11% and maximum of 100% while for the Wine dataframe an average of 87.96% is obtained, a minimum outlier value of 81.48% and a maximum outlier of 94.44%. The best distribution in the test phase is obtained for the Wine dataframe despite its values are significantly lower than those obtained by the Iris dataframe.

Using the Euclidean metric, the MDC performs better both in the train phase and in the test phase for the Iris dataframe.

MDC Cityblock Performance EVALUATION



Iris dataframe

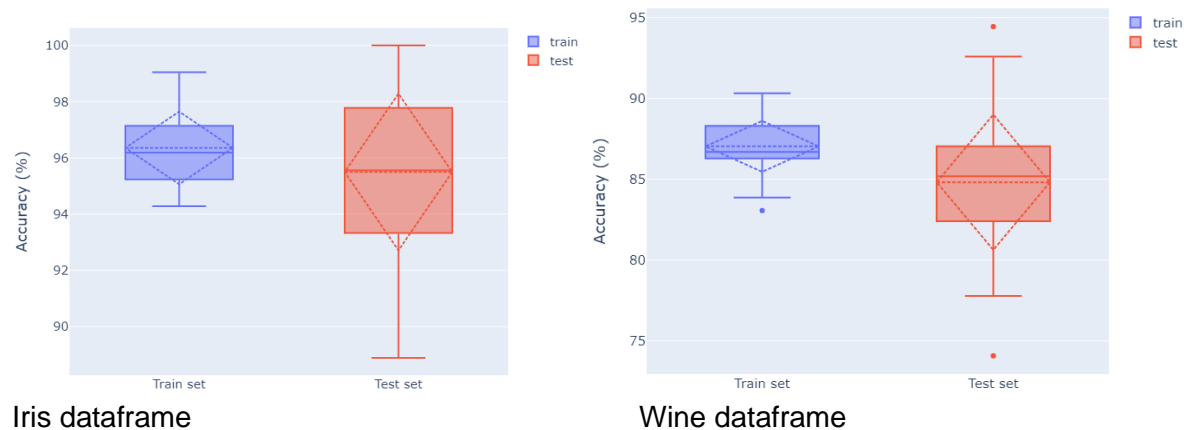


Wine dataframe

As you can see from the boxplots, the MDC with Cityblock metric, performs better in terms of accuracy, in the test phase, on the Iris dataframe, with an average accuracy of 95.39%, minimum value of 88.89% and maximum of 100% while for the Wine dataframe an average of 88.80% is obtained, a minimum value of 77.78% and a maximum of 96.29%. Also in the test phase there is a better distribution of the data for the Iris dataframe.

With the Cityblock metric, the MDC performs better both in the train phase and in the test phase for the Iris dataframe.

MDC Chessboard Performance EVALUATION




According to the boxplots, the MDC with Chessboard metrics performs better in terms of accuracy, in the test phase, on the Iris dataframe, with an average accuracy of 95.50%, a minimum value of 88.89% and a maximum of 100% while for the Wine dataframe an average of 84.81% is obtained, a minimum outlier value of 74.07% and a maximum outlier of 94.44%. With this metric you have a better distribution in the test phase for the Iris dataframe.

Using the Chessboard metric, the MDC performs better in both the train and test phase for the Iris dataframe.

Considering the results seen so far we can say that the MDC achieves better accuracy values with all three metrics for the Iris dataframe. The distribution of these values, most of the time, is better for the Iris dataframe while in the case with Euclidean metrics the distribution of the data is in favor of the Wine dataframe, despite having obtained worse accuracy values.

KNN Performance EVALUATION (Iris dataframe)



The analysis continues with the variation of the k value in the test phase, where we noticed a rapid increase in accuracy right from the start (k = 1 93%, k = 4 100%) with the increase of the k value by a few unit. The accuracy is kept stable up to a high number of k (k = 25) and then decreases in value (k = 26-30 95.5%) and then rises


swinging up to stabilize with very high ks (k = 35- 40 100%). The accuracy in these 40 k remains high in most cases.

```
[[18 0 0]
 [ 0 11 1]
 [ 0 2 13]]
```

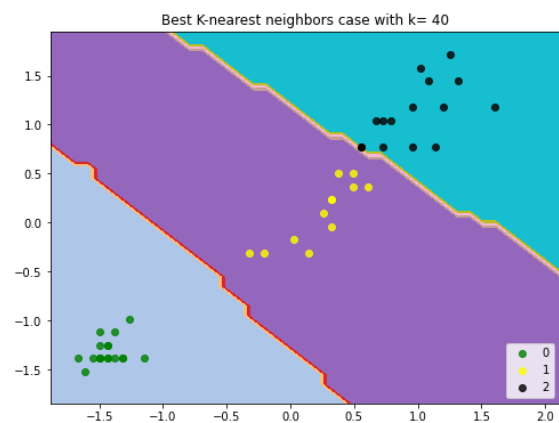
	precision	recall	f1-score	support
0	1.00	1.00	1.00	18
1	0.85	0.92	0.88	12
2	0.93	0.87	0.90	15
accuracy			0.93	45
macro avg	0.92	0.93	0.93	45
weighted avg	0.94	0.93	0.93	45

```
[[18 0 0]
 [ 0 12 0]
 [ 0 0 15]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	18
1	1.00	1.00	1.00	12
2	1.00	1.00	1.00	15
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

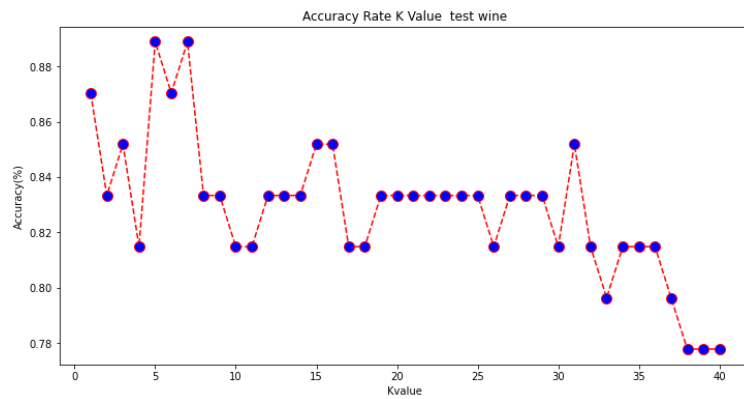


Worst k for Iris dataframe (k=1)



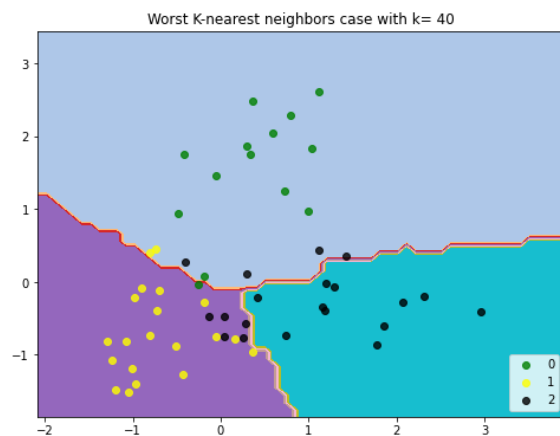
One of the best k for Iris dataframe (k=40)

KNN Performance EVALUATION (Wine dataframe)




By varying the value of k in the test phase for the Wine dataframe, we notice a fluctuating trend, with 2 points of maximum accuracy ($k = 5$ and $k = 7$, 89%). The accuracy level in the general tendency is decreasing until reaching very high k values ($k = 38, 39, 40$) where the lowest accuracy level is obtained (78%).

[[13 1 0] [2 18 0] [4 5 11]]					[[14 0 0] [1 16 3] [2 0 18]]				
	precision	recall	f1-score	support		precision	recall	f1-score	support
0	0.68	0.93	0.79	14	0	0.82	1.00	0.90	14
1	0.75	0.90	0.82	20	1	1.00	0.80	0.89	20
2	1.00	0.55	0.71	20	2	0.86	0.90	0.88	20
accuracy			0.78	54	accuracy			0.89	54
macro avg	0.81	0.79	0.77	54	macro avg	0.89	0.90	0.89	54
weighted avg	0.83	0.78	0.77	54	weighted avg	0.90	0.89	0.89	54



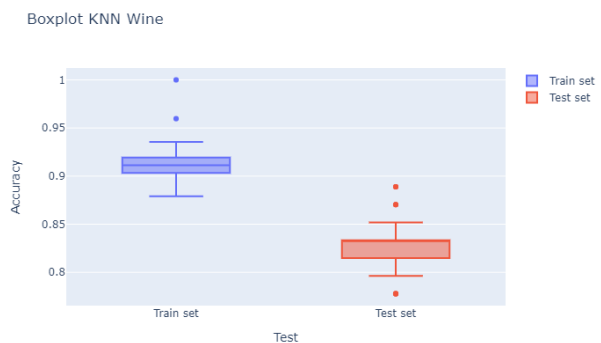
One of the worst k for Wine dataframe ($k=40$)



One of the best k for Wine dataframe ($k=5$)



Iris dataframe KNN boxplots




Wine dataframe KNN boxplots

The distribution of the accuracy values in the test phase is better for the Iris dataframe, with a range that varies from 0.98 to 1. There are also 2 outliers of which the minimum is 0.93. Regarding the Wine dataframe, its values are in a range that varies between 0.79 and 0.85. There are 3 outliers of which a minimum with a value of 0.78 and a maximum of 0.89. The KNN in the test phase performs better on the Iris dataframe than the Wine dataframe, which performed worse than the training phase.

Considering the results obtained, we can affirm that the KNN performs better both in the training phase and in the test phase with the Iris dataframe.

FINAL RESULTS



According to the results obtained, it is clear that the MDC performs better than the KNN.

In terms of accuracy, the best result was obtained with the MDC with Euclidean metric on the Iris dataframe, while the worst was obtained with the MDC with Chessboard metric on the Wine dataframe.

By averaging the results obtained by varying the distance metric on the MDC, the best result is obtained with the MDC with Euclidean metric and the worst is obtained with the MDC with Chessboard metric.

Based on the difference in the average accuracy value on the two dataframes, the classifier that gave us the closest accuracy values for both data frames is the MDC with Cityblock metric and the one with the most different values is the KNN.

In our opinion, the MDC with Cityblock metric performs better as it has managed to provide more similar accuracy values than all the other metrics used for the MDC and KNN.

Bibliography

- Slides from Intelligent System course
- Slides from Data Science for Business course
- scikit-learn.org
- pandas.pydata.org
- plotly.com
- medium.com